


# IMD0033 - Probabilidade

## Aula 09 - Introdução a Pandas I

Ivanovitch Silva  
Setembro, 2018



# Agenda

---

- About the two core pandas types: dataframes and series
- How to select data using row and column labels
- A variety of methods for exploring data with pandas
- How to assign data using various techniques in pandas
- How to use boolean indexing with pandas for selection and assignment

# Atualizar o repositório

---

```
git clone https://github.com/ivanovitchm/imd0033_2018_2.git
```

Ou ....

```
git pull
```

# Understanding Pandas & Numpy

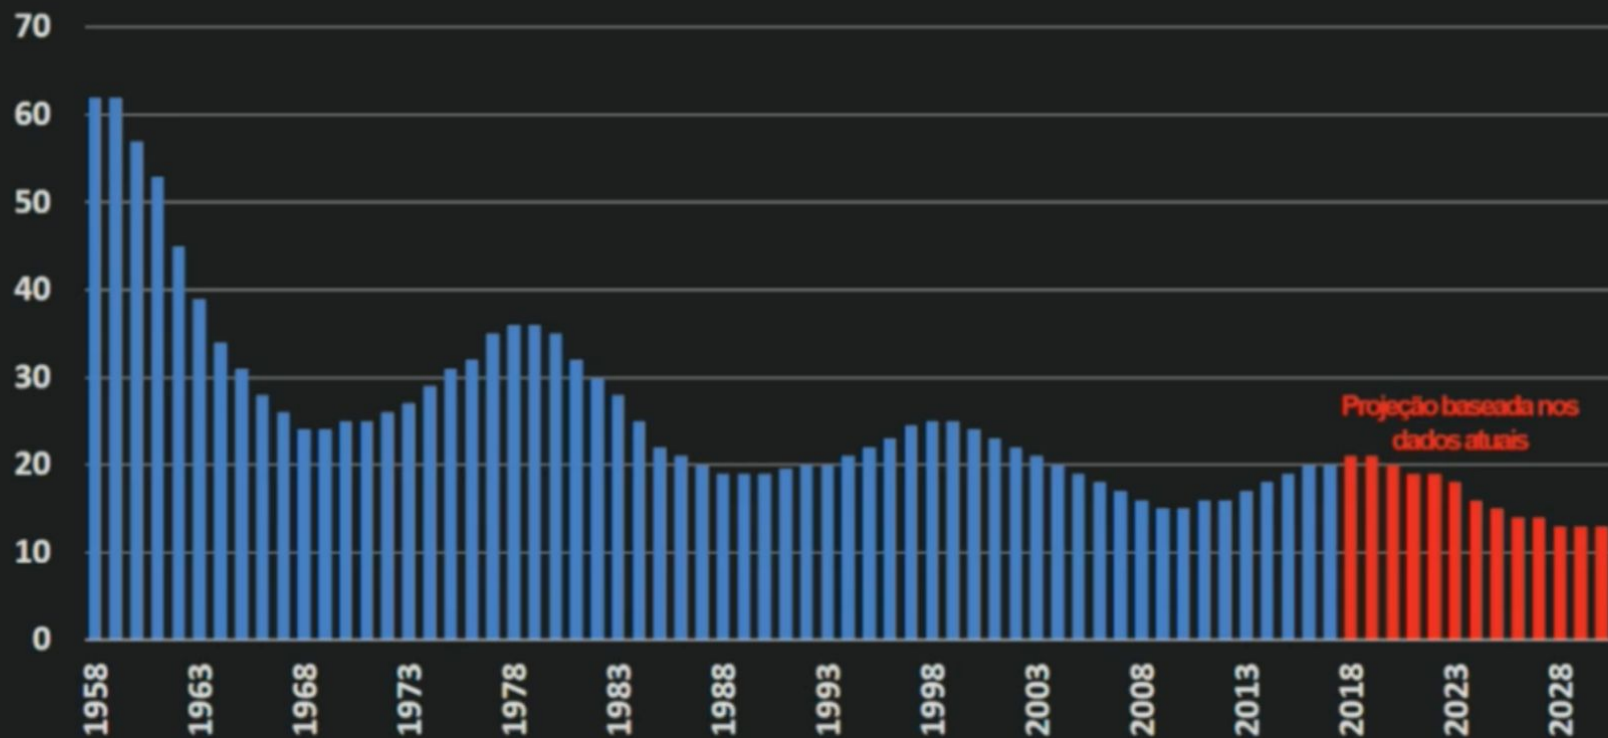
---

- Numpy
  - a. Lack support for column names
  - b. Support for only one data type per ndarray
  - c. There are lots of low level methods, however there are many common analysis patterns that don't have pre-built methods.

**The pandas library** provides solutions to all of these pain points and more. Pandas is not so much a replacement for NumPy as an **extension of NumPy**.



## Tempo médio de permanência de uma empresa no S&P 500 (em anos)



Fonte: INNOSIGHT, Richard N. Foster, Standard & Poor's

# UMA EMPRESA DO S&P 500 ESTÁ SENDO SUBSTITUÍDA A CADA DUAS SEMANAS

Richard Foster


# The dataset

	rank	revenues	revenue_change	profits	assets	profit_change	ceo	industry	sector	previous_rank
<b>Walmart</b>	1	485873	0.8	13643.0	198825	-7.2	C. Douglas McMillon	General Merchandisers	Retailing	1
<b>State Grid</b>	2	315199	-4.4	9571.3	489838	-6.2	Kou Wei	Utilities	Energy	2
<b>Sinopec Group</b>	3	267518	-9.1	1257.9	310726	-65.0	Wang Yupu	Petroleum Refining	Energy	4
<b>China National Petroleum</b>	4	262573	-12.3	1867.5	585619	-73.7	Zhang Jianhua	Petroleum Refining	Energy	3
<b>Toyota Motor</b>	5	254694	7.7	16899.3	437575	-12.3	Akio Toyoda	Motor Vehicles and Parts	Motor Vehicles & Parts	8

```
import pandas as pd
f500 = pd.read_csv("f500.csv", index_col=0)
f500.index.name = None
```



# Introducing Dataframes



# Introducing Dataframes

---

```
# put your code here  
f500.info(,)
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 500 entries, Walmart to AutoNation  
Data columns (total 16 columns):  
rank                500 non-null int64  
revenues            500 non-null int64  
revenue_change      498 non-null float64  
profits             499 non-null float64  
assets              500 non-null int64  
profit_change       436 non-null float64  
ceo                  500 non-null object  
industry             500 non-null object  
sector              500 non-null object  
previous_rank        500 non-null int64  
country              500 non-null object  
hq_location          500 non-null object  
website              500 non-null object  
years_on_global_500_list 500 non-null int64  
employees            500 non-null int64  
total_stockholder_equity 500 non-null int64  
dtypes: float64(3), int64(7), object(6)  
memory usage: 66.4+ KB
```

f500.head()  
f500.tail()

# Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
f500_selection	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

```
f500_selection.loc[:, "rank"]
```

Walmart	1
State Grid	2
Sinopec Group	3
China Natural Petroleum	4
Toyota Motor	5

# Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
f500_selection	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

		country	rank
f500_selection.loc[:, ["country", "rank"]]	Walmart	USA	1
	State Grid	China	2
	Sinopec Group	China	3
	China Natural Petroleum	China	4
	Toyota Motor	Japan	5

# Selecting Columns From a Dataframe by label

		rank	revenues	profits	country
f500_selection	Walmart	1	485873	13643.0	USA
	State Grid	2	315199	9571.3	China
	Sinopec Group	3	267518	1257.9	China
	China Natural Petroleum	4	262573	1867.5	China
	Toyota Motor	5	254694	16899.3	Japan

		rank	revenues	profits
f500_selection.loc[:, "rank": "profits"]	Walmart	1	485873	13643.0
	State Grid	2	315199	9571.3
	Sinopec Group	3	267518	1257.9
	China Natural Petroleum	4	262573	1867.5
	Toyota Motor	5	254694	16899.3

# Column selection shortcuts

---


Select by Label	Explicit Syntax	Common Shorthand	Other Shorthand
Single column	<code>df.loc[:, "col1"]</code>	<code>df["col1"]</code>	<code>df.col1</code>
List of columns	<code>df.loc[:, ["col1", "col7"]]</code>	<code>df[["col1", "col7"]]</code>	
Slice of columns	<code>df.loc[:, "col1":"col4"]</code>		

# Lesson\_09\_Introduction\_to\_pandas.ipynb

## Up to section 1.4



# Selecting Items from a Series by Label





Original  
Dataframe

	A	B	C	D	E
V					
W					
X					
Y					
Z					

## Code

```
single_col = df["D"]
```

## Result

V	
W	
X	
Y	
Z	

**single\_col** is a  
series object

	A	B	C	D	E
V					
W					
X					
Y					
Z					

```
single_row = df.head(1)
```

A	
B	
C	
D	
E	

**single\_row** is a  
series object

Original  
Dataframe

	A	B	C	D	E
V					
W					
X					
Y					
Z					

## Code

```
multi_cols = df[["A", "C", "D"]]
```

## Result

	A	C	D
V			
W			
X			
Y			
Z			

**multi\_cols** is a  
dataframe object

	A	B	C	D	E
V					
W					
X					
Y					
Z					

```
multi_rows = df.head(3)
```

	A	B	C	D	E
V					
W					
X					

**multi\_rows** is a  
dataframe object

# Dataframe vs Series

---

	<b>Series</b>	<b>DataFrame</b>
<b>Dimensions</b>	One	Two
<b>Has 'index' axis</b>	Yes	Yes
<b>Has 'columns' axis</b>	No	Yes
<b>Number of dtypes</b>	One	Many (one per column)

# Series and Dataframe Describe Methods

---

```
revs = f500["revenues"]  
print(revs.describe())
```

```
count      500.000000  
mean      55416.358000  
std       45725.478963  
min       21609.000000  
25%       29003.000000  
50%       40236.000000  
75%       63926.750000  
max       485873.000000  
Name: revenues, dtype: float64
```

```
print(f500["assets"].describe())
```

```
count      5.000000e+02  
mean      2.436323e+05  
std       4.851937e+05  
min       3.717000e+03  
25%       3.658850e+04  
50%       7.326150e+04  
75%       1.805640e+05  
max       3.473238e+06  
Name: assets, dtype: float64
```


# Lesson\_09\_Introduction\_to\_pandas.ipynb

## Sections 1.5, 1.6, 1.7



`DataFrame.method(axis=0)`  
or  
`DataFrame.method(axis="index")`


Calculates along the **row** axis



Calculates result for  
each **column**.

`DataFrame.method(axis=1)`  
or  
`DataFrame.method(axis="column")`

Calculates along the **column** axis



Calculates result for  
each **row**.

## More data exploration methods

```
medians = f500[["revenues", "profits"]].median(axis=0)
# we could also use .median(axis="index")
print(medians)
```

```
revenues    40236.0
profits      1761.6
dtype: float64
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	485873
State Grid	2	315199
Sinopec Group	3	267518
China National Petroleum	4	262573
Toyota Motor	5	254694

```
>>> top5_rank_revenue["revenues"] = 0
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	0
State Grid	2	0
Sinopec Group	3	0
China National Petroleum	4	0
Toyota Motor	5	0

## Assignment with Pandas

# Assignment with Pandas

---

```
>>> top5_rank_revenue.loc["Sinopec Group", "revenues"] = 999
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues
Walmart	1	0
State Grid	2	0
Sinopec Group	3	999
China National Petroleum	4	0
Toyota Motor	5	0



# Add a new column

---

```
>>> top5_rank_revenue["year_founded"] = 0
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues	year_founded
Walmart	1	0	0
State Grid	2	0	0
Sinopec Group	3	999	0
China National Petroleum	4	0	0
Toyota Motor	5	0	0

# Add a new row

---

```
>>> top5_rank_revenue.loc["My New Company"] = 555
```

```
>>> print(top5_rank_revenue)
```

	rank	revenues	year_founded
Walmart	1	0	0
State Grid	2	0	0
Sinopec Group	3	999	0
China National Petroleum	4	0	0
Toyota Motor	5	0	0
My New Company	555	555	555

# Using boolean indexing with pandas objects

w	2
x	4
y	6
z	8

pandas series



w	2	< 5
x	4	< 5
y	6	< 5
z	8	< 5

Vectorized  
boolean operation



w	True
x	True
y	False
z	False

Boolean  
pandas series

	A	B
w	2	3
x	4	6
y	6	9
z	8	12

pandas DataFrame



	A	B	
w	2	3	< 5
x	4	6	< 5
y	6	9	< 5
z	8	12	< 5

Vectorized  
boolean operation



	A	B
w	True	True
x	True	False
y	False	False
z	False	False

Boolean  
pandas DataFrame

# Using boolean indexing with pandas objects

```
result = df.loc[num_bool, "name"]
```

		name		
False	w	Kylie	12	
True	→ x	Rahul	8	→ x
False	y	Michael	5	z
True	→ z	Sarah	8	↗
				<b>result</b>
				Rahul
				Sarah

```
result = df[num_bool]
```

		name	num	
False	w	Kylie	12	
True	→ x	Rahul	8	→ x
False	y	Michael	5	z
True	→ z	Sarah	8	↗
				<b>result</b>
				name num
				Rahul 8
				Sarah 8

# Using boolean arrays to assign values

---

```
f500.loc[f500["sector"] == "Motor Vehicles & Parts", "sector"] = "Motor Vehicles and Parts"
```

# Challenge

---

## Finding top performers by country

```
>>> top_3_countries = f500["country"].value_counts().head(3)
```

```
>>> print(top_3_countries)
```

```
USA      132
```

```
China    109
```

```
Japan     51
```

```
Name: country, dtype: int64
```

# Lesson\_09\_Introduction\_to\_pandas.ipynb

